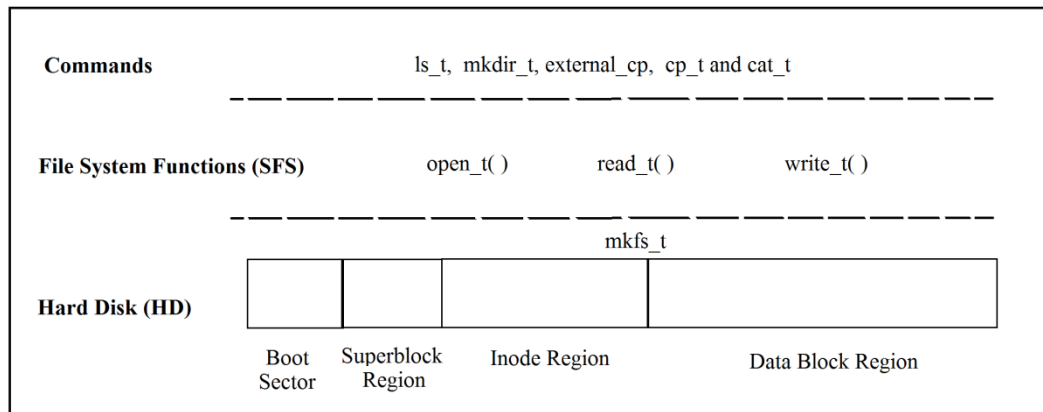


# CSCI 3150 Introduction to Operating Systems

## Assignment Three

Deadline: 23:55, 9 December 2018

In this assignment, you are required to implement a simple file system called *SFS*. After you finish this homework, you should have better understanding of a file system and its organization/implementation. An overview of SFS is shown in the figure below.

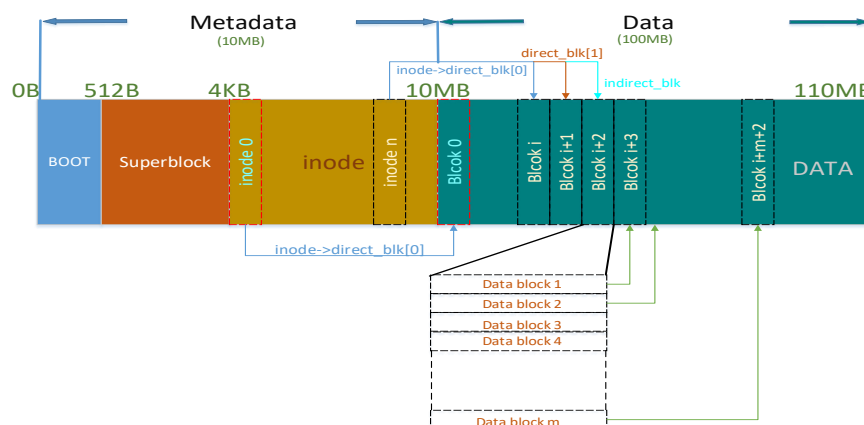


### 1. SFS (Simple File System)

SFS works on a file called **HD** that is a 110MB file (initially empty) and you can find it from the zip file provided.

The implementation of *SFS* consists two parts: three filesystem related functions (`open_t()`, `read_t()`, and `write_t()`), and six commands (`mkfs_t`, `ls_t`, `mkdir_t`, `external_cp`, `cp_t`, `cat_t`).

- Three file-system-related functions are based on the simple file system. An illustrative format on HD is shown below:



As shown above, in HD, there are two regions: the metadata and data regions. The metadata region is inside the first 10MB; it contains a boot sector (the first 512 bytes), the superblock and inode regions. The superblock region is from 512B to 4KB, and the inode

region from 4KB to 10MB. The data region is from 10 MB to 110 MB, in which it is divided into data blocks (each data block is 4 KB).

The superblock region defines the layout and its format can be found from the following structure:

```
struct  superblock      /*The key information of filesystem */
{
    int  inode_offset;    /* The start offset of the inode region */
    int  data_offset;     /* The start offset of the data region */
    int  max_inode;       /* The maximum number of inodes */
    int  max_data_blk;    /* The maximum number of data blocks */
    int  next_available_inode; /* The index of the next free inode */
    int  next_available_blk; /* The index of the next free block */
    int  blk_size;        /* The size per block */
};
```

Basically, the inode region starts at 4 KB (inode\_offset); the data region starts at 10 MB (data\_offset), the maximum number of inodes is 100 (max\_inode); the maximum number of data blocks is 25600; next\_available\_inode and next\_available\_blk are used to represent the indexes of the next free inode and the next free block, respectively; the block size is 4 KB. To make it simple, you do not need to reclaim inodes or data blocks, and you can simply obtain the next available inode (data block) index based on next\_available\_inode (next\_available\_blk) when you create a file (allocate data blocks).

The inode region contains inodes that can be retrieved based on its index in the inode region (called the inode number). An inode is used to represent a file, and is defined based on the following structure:

```
struct inode            /* The structure of inode, each file has only one inode */
{
    int  i_number;       /* The inode number */
    time_t i_mtime;      /* Creation time of inode */
    int  i_type;         /* Regular file for 0, directory file for 1 */
    int  i_size;          /* The size of file */
    int  i_blocks;        /* The total numbers of data blocks */
    int  direct_blk[2];   /* Two direct data block pointers */
    int  indirect_blk;    /* One indirect data block pointer */
    int  file_num;        /* Number of files under a directory (0 for regular file) */
};
```

Some related parameters can be found as follows:

```
#define SB_OFFSET    512          /* The offset of superblock region */
#define INODE_OFFSET  4096        /* The offset of inode region */
#define DATA_OFFSET  10485760    /* The offset of data region */
#define MAX_INODE     100         /* The maximum number of inode */
#define MAX_DATA_BLK  25600       /* The maximum number of block */
#define BLOCK_SIZE    4096        /* The size per block */
#define MAX_NESTING_DIR 10         /* The nesting number of directory */
#define MAX_COMMAND_LENGTH 50    /* The maximum command length */
```

In SFS, an inode contains two direct data block pointers and one single indirect data block pointer. There are two types of files: regular and directory files. The **content of a directory file** should follow the following structure:

```
typedef struct dir_mapping /* Record file information in directory file */
{
    char  dir[20];          /* The file name in current directory */
    int   inode_number;     /* The corresponding inode number */
}DIR_NODE;
```

Each directory file should at least contain two mapping items, “.” and “..”, for itself and its parent directory, respectively.

Based on SFS, the prototypes of the three filesystem-related functions are shown as follows:

- 1) `int open_t(const char *pathname, int flags);`

Description: Given an absolute *pathname* for a file, `open_t()` returns the corresponding inode number of the file or **-1 if an error occurs**. The returned inode number will be used in subsequent functions in `read_t()` and `write_t()`.

The argument *flags* can be one of the following three values: **0 (or 1) means that a new regular (or directory) file will be created** (if one file with the **same name exists**, the new file will **replace the old file**); **2 means that the target is an existing file**. The new inode number can be obtained from *next\_available\_inode* from the superblock; after it is obtained, *next\_available\_inode* should be incremented by one.

- 2) `int read_t(int inode_number, int offset, void *buf, int count);`

Description: `read_t()` attempts to **read up to *count* bytes** from the file **starting at *offset*** (with the inode number *inode\_number*) into the buffer starting at *buf*. It commences at the file offset specified by *offset*. **If *offset* is at or past the end of file, no bytes are read**, and `read_t()` returns zero. On success, the **number of bytes read is returned** (zero indicates end of file), and on error, -1 is returned.

- 3) `int write_t(int inode_number, int offset, void *buf, int count);`

Description: `write_t()` writes up to *count* bytes from the buffer pointed *buf* to the file referred to by the inode number *inode\_number* starting at the file offset at *offset*. The number of bytes written may be less than *count* if there is insufficient space on the underlying physical medium or the maximum size of a file has been achieved. On success, the number of bytes written is returned (zero indicates nothing was written). On error, -1 is returned.

A new block can be obtained from *next\_available\_blk*; after it is obtained, *next\_available\_blk* should be incremented by one.

- Six commands: `mkfs_t`, `ls_t`, `mkdir_t`, `external_cp`, `cp_t`, and `cat_t`. Based on the above, six commands are used to support SFS. Among them, `mkfs_t` directly works on HD, and the other commands (`ls_t`, `mkdir_t`, `external_cp`, `cp_t`, and `cat_t`) work with SFS based on the above three functions (`open_t()`, `read_()`, and `write()`). They are described as follows:

1) `mkfs_t`

Description: `mkfs` is used to build an SFX filesystem on HD. **This should be the first step in order to use our SFS filesystem on HD.** After this command is successfully executed, the parameters in the superblock region discussed above will be set up correspondingly in HD. Moreover, the first inode (whose inode number is 0) is created for the root directory, and its initial data (related to “.” and “..”) will be written into Block 0 based on `DIR_NODE`. See Lab 9 for details.

2) `ls_t dname`

Description: `ls_t` lists the information of all files under the directory *dname* and *dname* should be an absolute pathname. The information related to each file should include its inode number, creation time, file type (regular or directory), and the size of the file.

3) `mkdir_t dname`

Description: `mkdir_t` creates a new directory file with the name *dname* and *dname* includes the new directory name with the absolute path. The new directory file will be created even if this is a directory file with the same name (i.e. the new directory file will replace the old one under the directory).

4) `external_cp outside_path_name sfs_path_name`

Description: `external_cp` copies a regular file from Linux (specified by *outside\_path\_name* that is the absolute path) to a file (with *sfs\_path\_name* as the absolute path and name) inside the SFS filesystem. A new regular file will be created and copied in the SFS if the path/names specified by *outside\_path\_name* and *sfs\_path\_name* are effective (the new regular file will be created and copied in the SFS even if there is a regular file with the same path/name *sfs\_path\_name*); otherwise, the error will be reported.

5) `cp_t source_path_name destination_path_name`

Description: `cp_t` copies a regular file (specified by *source\_path\_name* that is the absolute path) to the destination (specified by *destination\_path\_name* that is the absolute path). A new regular file will be created and copied in the SFS if the path/names specified by *source\_path\_name* and *destination\_path\_name* are effective (an old file with the same path/name as *destination\_path\_name* will be replaced by the new file); otherwise, the error will be reported.

6) `cat_t path_name`

Description: `cat_t` prints the contents of the file specified by the absolute path/name *file\_name* to the standard output. If the file does not exist, the error will be reported.

## 2. Requirements

In this assignment, you need to implement `open_t()`, `read_t()` and `write_t()`, thereby making SFS work with the above five commands (`ls_t`, `mkdir_t`, `external_cp`, `cp_t`, and `cat_t`). Note that the source code for `mkfs_t` can be found in Lab 9, and the source code for the other five commands has been provided but `open_t()`, `read_t()` and `write_t()` have not been implemented.

To start, please download `assign3-code.zip`. After unzipping this zip file, you can find the following files:

- *call.c*: The source code for `open_t()`, `read_t()` and `write_t()` that you should implement. In *call.c*, you are allowed to create any auxiliary functions that can help your implementation. But only “`open_t()`”, “`wrtie_t()`”, and “`read_t()`” are allowed to call these auxiliary functions.
- *cat\_t.c*, *cp\_t.c*, *external\_cp.c*, *ls\_t.c*, *mkdir\_t.c*: The source code for the five commands mentioned above (`cat_t`, `cp_t`, `external_cp`, `ls_t`, and `mkdir_t`) that have been implemented.
- *call.h*, *inode.h*, *superblock.h*: The header files that define the data structures and function prototypes.
- *HD*: The hard disk file (110 MB);
- *Makefile*: The makefile file.

You can compile and generate all commands by inputting:

```
make
```

“make” will work based on “Makefile”. Note that HD is empty in the beginning; so the first step is to format HD with “`mkfs_t`” (see Lab 9 for details).

This assignment will be graded by Mr. Lei Zhu (Email: [leizhu@cse.cuhk.edu.hk](mailto:leizhu@cse.cuhk.edu.hk)). Your programs **must be able to be compiled/run under the XUbuntu environment (in Lab One)**.

**What to submit – A zip file that ONLY contains *call.c*.**